



# ClickHouse at MessageBird

Analysing billions of events in real-time\*

Aleksandar Aleksandrov & Félix Mattrat

NOVEMBER 2018

# About us

Data engineers & Team leads



**Aleksandar  
Aleksandrov**



**Félix  
Mattrat**

ABOUT

# Introducing MessageBird

MessageBird is a cloud communications platform that empowers consumers to communicate with your business in the same way they communicate with their friends - seamlessly, on their own timeline and with the context of previous conversations.

For additional information visit: [www.messagebird.com](http://www.messagebird.com)

## **225+** Agreements

We have 225+ direct-to-carrier agreements with operators worldwide.

## **15,000+** Customers

Customers in over 60+ countries, across a great variety of industries.

## **180+** Employees

More than 180 employees speaking over 20 languages based in the Americas, Europe & Asia.



# Our offices

Our bird's nests around the world provide customers  
in every time zone with 24/7 support.




AMSTERDAM



SINGAPORE




HAMBURG



LONDON



SAN FRANCISCO



SHANGHAI



SYDNEY



# What's on the menu?

- 01. Data at MessageBird
- 02. The past - Age Of Darkness
- 03. Enlightenment - ClickHouse use case
- 04. What's next? - Nirvana



# Needs

Mostly about statistics and reporting

## Internal needs

- State of the system
- Routing SMS
- Training algorithms
- ML Models

## External needs

- Customer dashboard
- Reporting API

# The landscape

- Multiple carriers is messy - no uniformity of the data
- SMS messages go through many state changes up to months into the past
- Pricing (both carrier and customer) changes retro-actively



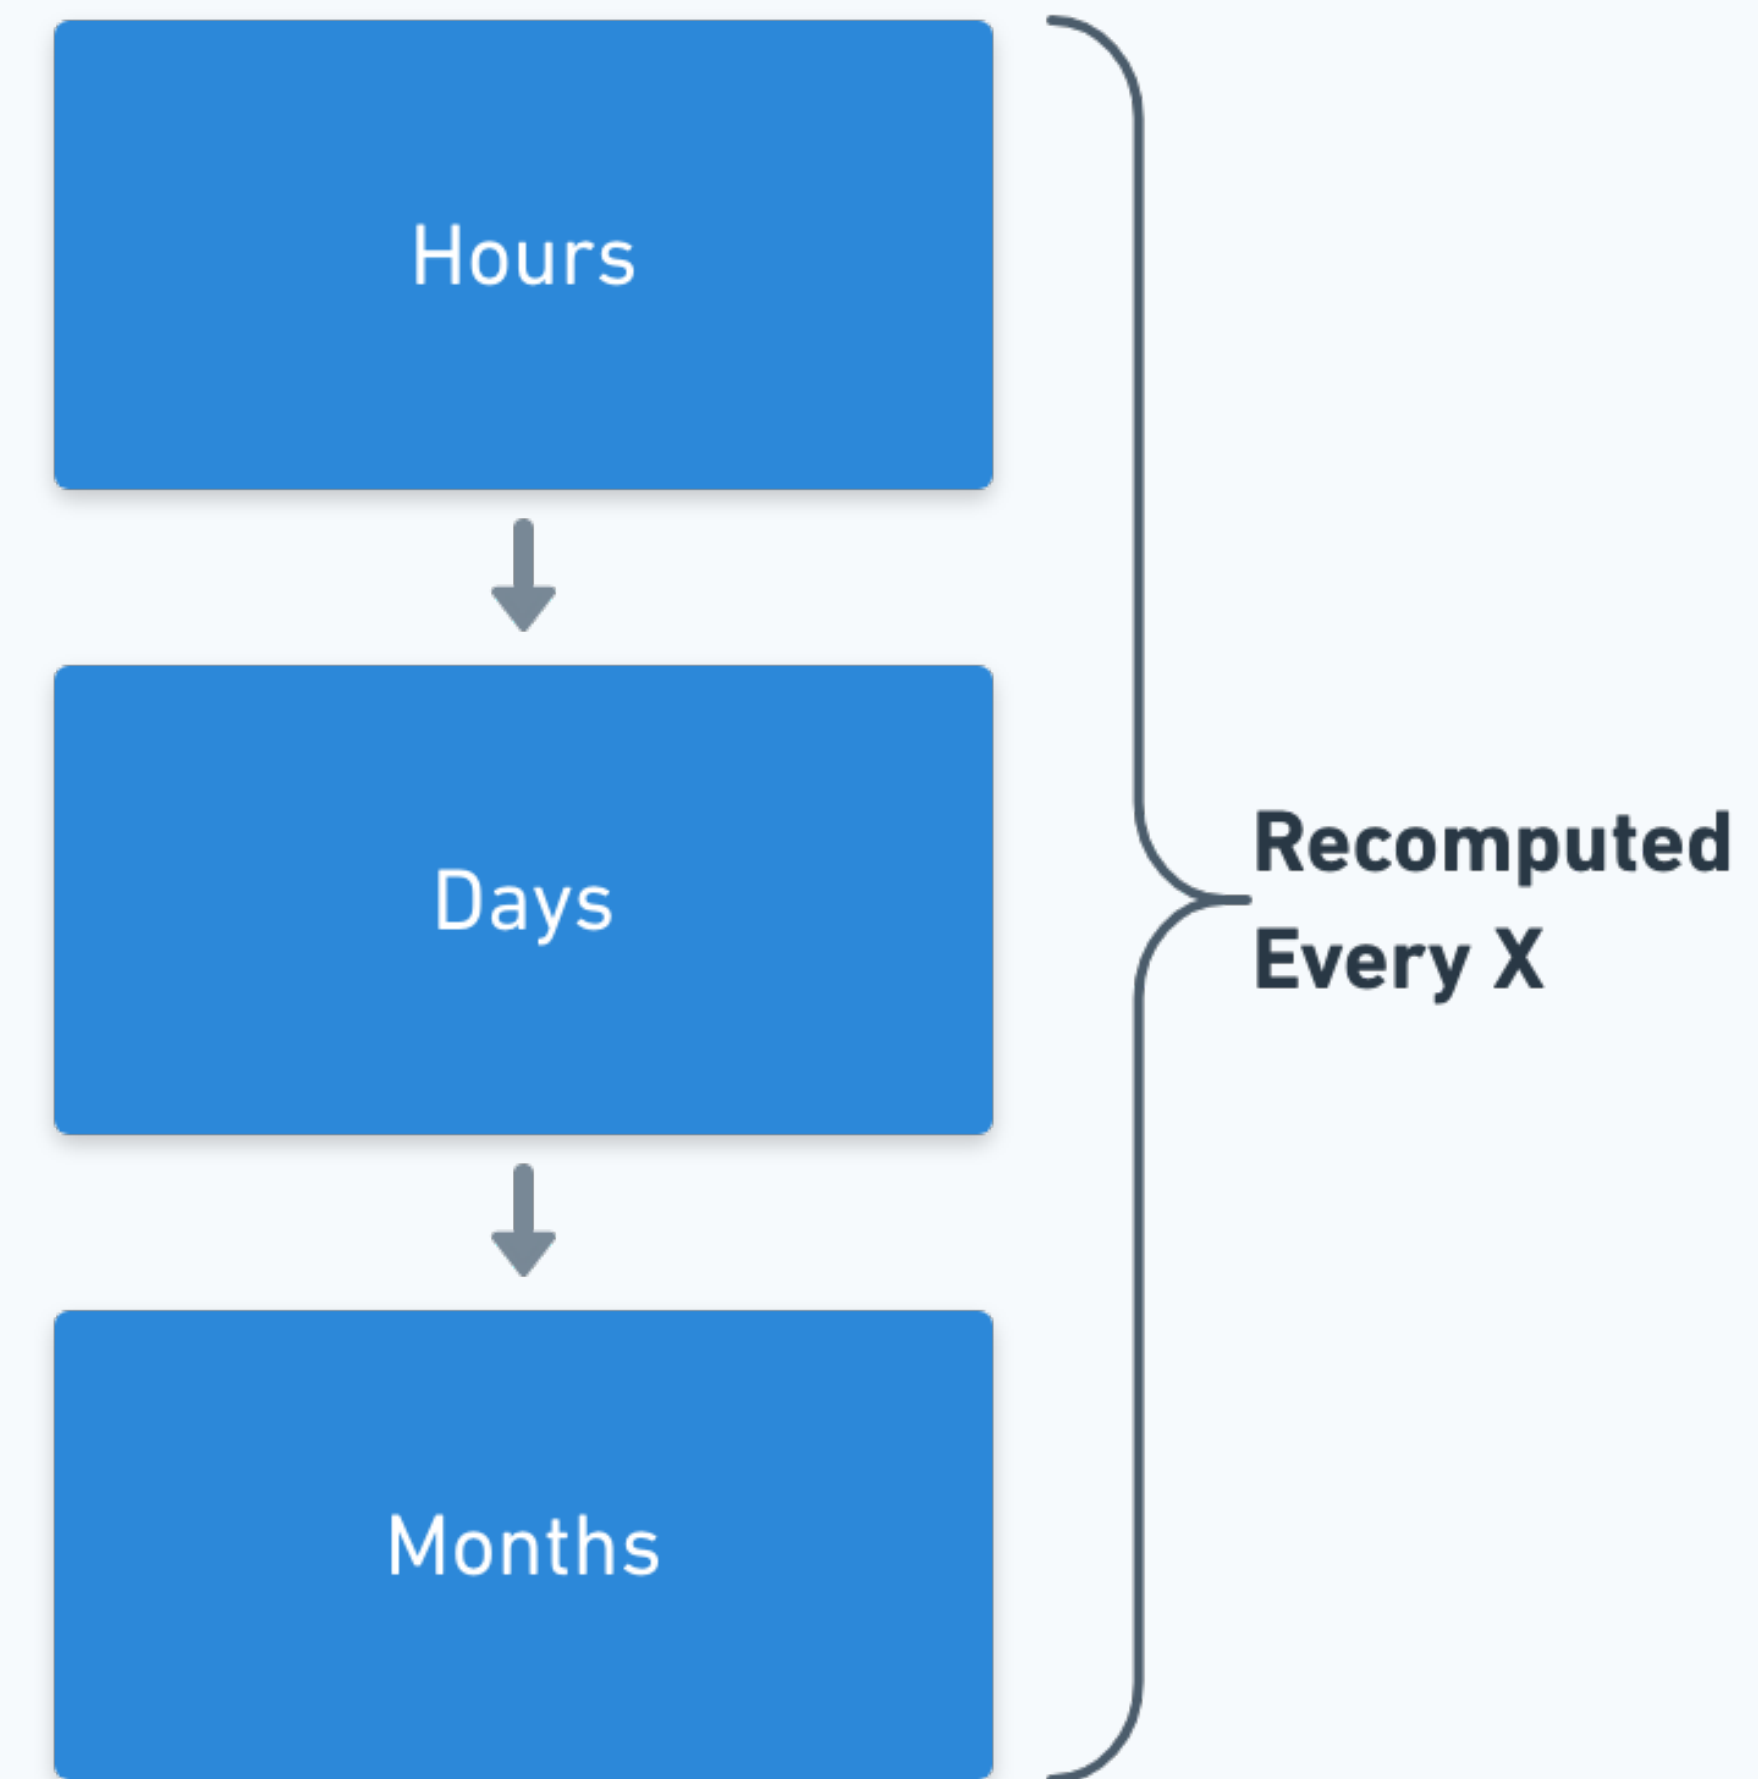
# Age of Darkness



AGE OF DARKNESS

# Hello CRON my old friend

- MySQL based
- Aggregates re-computed every X period of time
- Served us well for +5 years



AGE OF DARKNESS

# Scaling problems

- The system had difficulty scaling and was often lagging
- Loss of granularity with pre-aggregation
- Performed poorly while doing analytical queries
- Inaccuracies

ENLIGHTENMENT - QUEST FOR AN ALTERNATIVE

# Re-thinking data collection

- Able to keep up with continuously changing SMS message states
- In real time\*
- Scalable to handle MessageBird's global growth
- More flexible to accommodate wider use of data

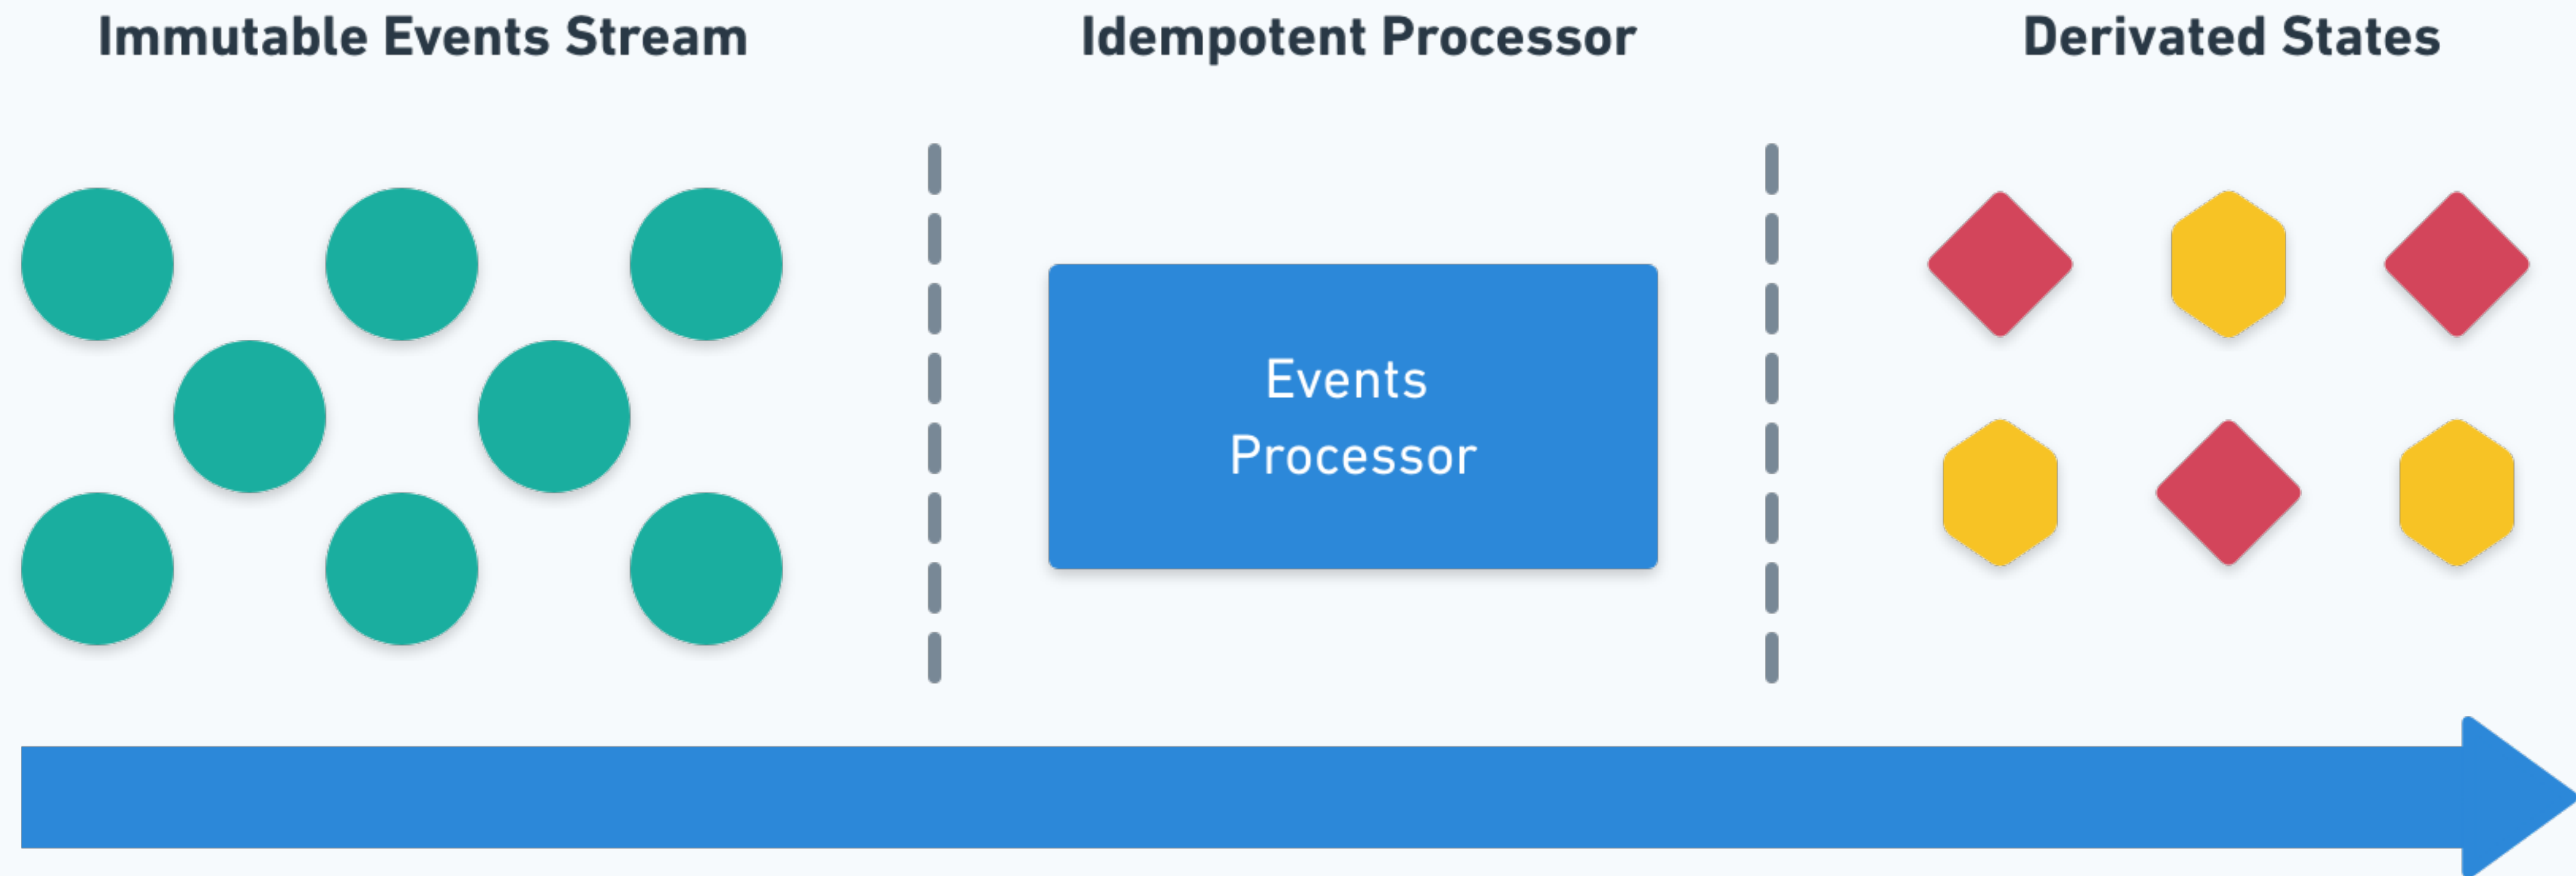


ENLIGHTENMENT - QUEST FOR AN ALTERNATIVE

# Introducing event sourcing

- Event sourcing, fairly common technique
- An immutable stream of events from which all states can be derivate



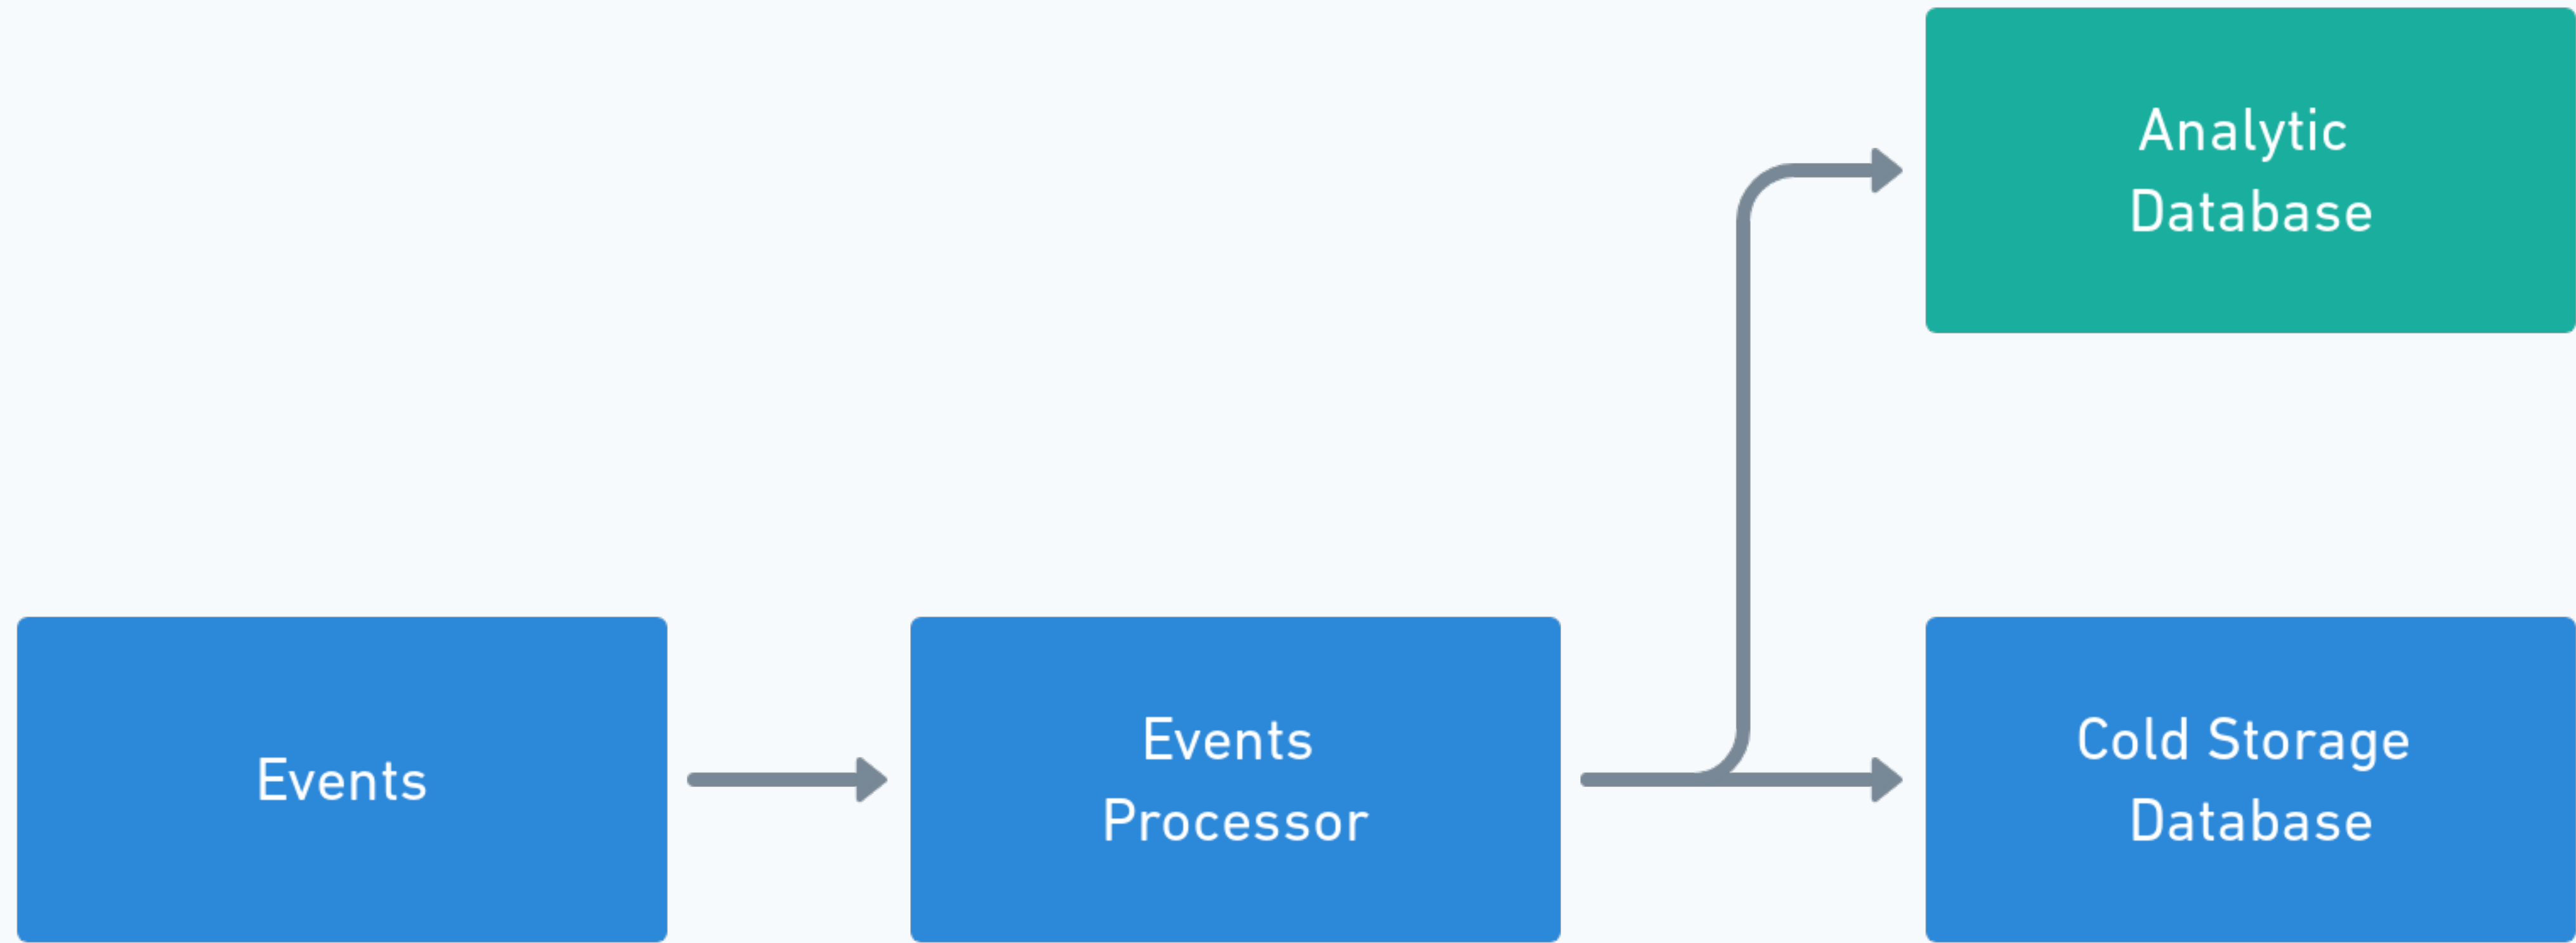


ENLIGHTENMENT - QUEST FOR AN ALTERNATIVE

# Introducing event sourcing

- Problem: now we have increased our data by an order of magnitude.
- How can we query this efficiently?





ENLIGHTENMENT - QUEST FOR AN ALTERNATIVE

# What is our unicorn database?

- Able to ingest large amount of data
- Data available immediately after ingestion
- No loss of granularity
- Flexible querying capabilities
- Sub-second response time
- Horizontally scalable





ENLIGHTENMENT - QUEST FOR AN ALTERNATIVE

# Vitess



- Let's shard the data
- Now we have N shards of problems
- Still has the limitations of MySQL
- Poor analytical support (at the time)

ENLIGHTENMENT - QUEST FOR AN ALTERNATIVE

# Kudu/Impala



- Promising, very clean and well defined SQL interface
- Compatible with HDFS & Parquets
- Column oriented
- But unable to reach sub-second querying time over billions of rows



ENLIGHTENMENT - QUEST FOR AN ALTERNATIVE

# Google BigQuery



- Scale well, millions or billions doesn't matter
- Fully managed: it's someone else problem
- Standard SQL support
- Not open source
- Not made for sub-second querying

# ClickHouse

February 15th, 2017



**Aleksandar 'Reasonable' Aleksandrov** 9:43 AM

<http://tech.marksblogg.com/billion-nyc-taxi-clickhouse.html>



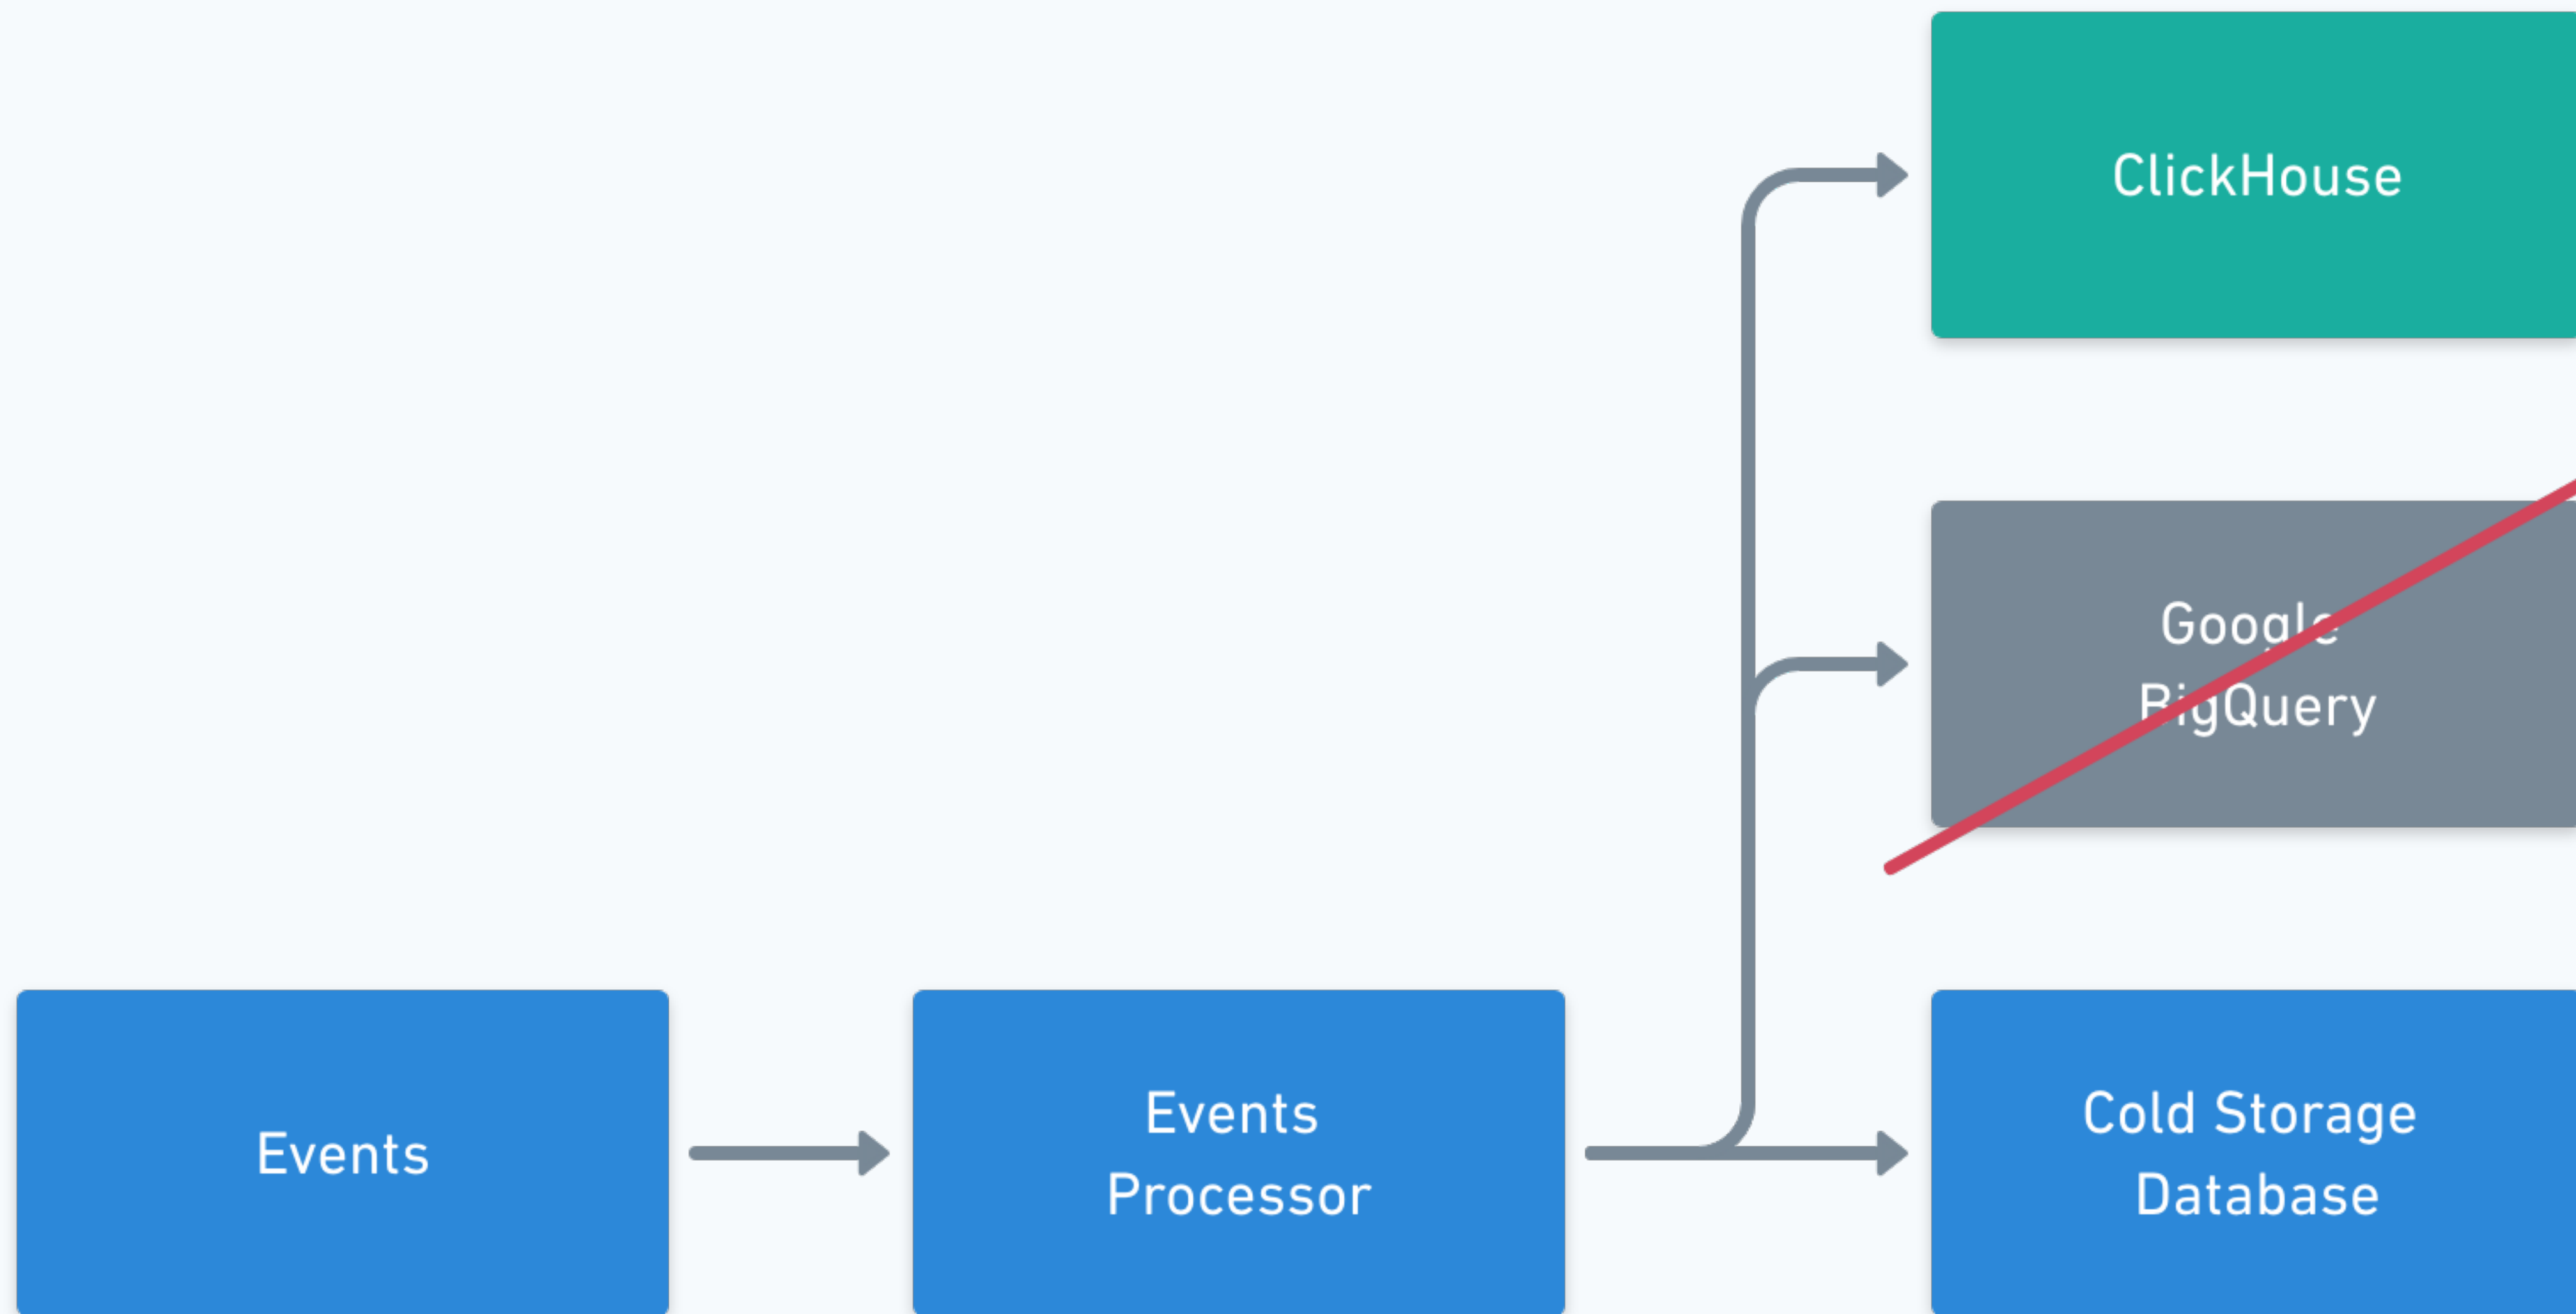
tech.marksblogg.com

## 1.1 Billion Taxi Rides on ClickHouse & an Intel Core i5

Benchmarks & Tips for Big Data, Hadoop, AWS, Google Cloud, Postgres, Spark, Python & More...

that looks pretty good





# ClickHouse



- Able to ingest a huge amount of data
- Sub-second on large dataset of non-aggregated data
- Flexible query capabilities: SQLish dialect
- Column oriented
- Scales **very** well vertically
- Horizontally scalable
- Open source

# ClickHouse

```
SELECT
    toStartOfQuarter(created_at) AS Quarter,
    mcc                          AS Country,
    floor(sum(sign * rate))      AS Total,
    sum(sign)                    AS MessageCount
FROM messages
WHERE created_at >= '2018-01-01' AND customer = 666
GROUP BY Quarter, Country
```

30 rows in set.

Elapsed: **0.33sec.**

Processed 497.91 million rows,  
4.95 GB

(1.42 billions rows/s., 14.39 GB/s.)

# ClickHouse what's the trick?

- Column oriented, you only pay for what you select
- Each column can potentially be processed in parallel
- Carefully crafted code makes use of vectorisation instructions
- Different table engines fit for different needs
- Horizontally scalable





**So, how to ingest ever changing data  
into ClickHouse**

# **CollapsingMergeTree**

- You write twice the amount of data, but eventually end up with a single row per PK
- Based on the idea of log compaction
- Excels at analytical queries on a large amount of data

# Collapsing what?

Primary key style

| sign | <u>date</u> | <u>id</u> | status   | price |
|------|-------------|-----------|----------|-------|
| 1    | 2018-10-08  | 666       | ACCEPTED | 0.01  |

# Collapsing what?

| sign | <u>date</u> | <u>id</u> | status   | price |
|------|-------------|-----------|----------|-------|
| 1    | 2018-10-08  | 666       | ACCEPTED | 0.01  |
| -1   | 2018-10-08  | 666       | ACCEPTED | 0.01  |

# Collapsing what?

| sign | <u>date</u> | <u>id</u> | status    | price |
|------|-------------|-----------|-----------|-------|
| 1    | 2018-10-08  | 666       | ACCEPTED  | 0.01  |
| -1   | 2018-10-08  | 666       | ACCEPTED  | 0.01  |
| 1    | 2018-10-08  | 666       | DELIVERED | 0.05  |

# Collapsing what?

```
SELECT sum(sign * price) AS total FROM dataset
```

| sign | <u>date</u> | <u>id</u> | status    | price |
|------|-------------|-----------|-----------|-------|
| 1    | 2018-10-08  | 666       | ACCEPTED  | 0.01  |
| -1   | 2018-10-08  | 666       | ACCEPTED  | 0.01  |
| 1    | 2018-10-08  | 666       | DELIVERED | 0.05  |

# Collapsing what?

```
SELECT sum(sign * price) AS total FROM dataset
```



# Collapsing what?

| sign          | <u>date</u>           | <u>id</u>      | status              | price           |
|---------------|-----------------------|----------------|---------------------|-----------------|
| <del>1</del>  | <del>2018-10-08</del> | <del>666</del> | <del>ACCEPTED</del> | <del>0.01</del> |
| <del>-1</del> | <del>2018-10-08</del> | <del>666</del> | <del>ACCEPTED</del> | <del>0.01</del> |
| 1             | 2018-10-08            | 666            | DELIVERED           | 0.05            |

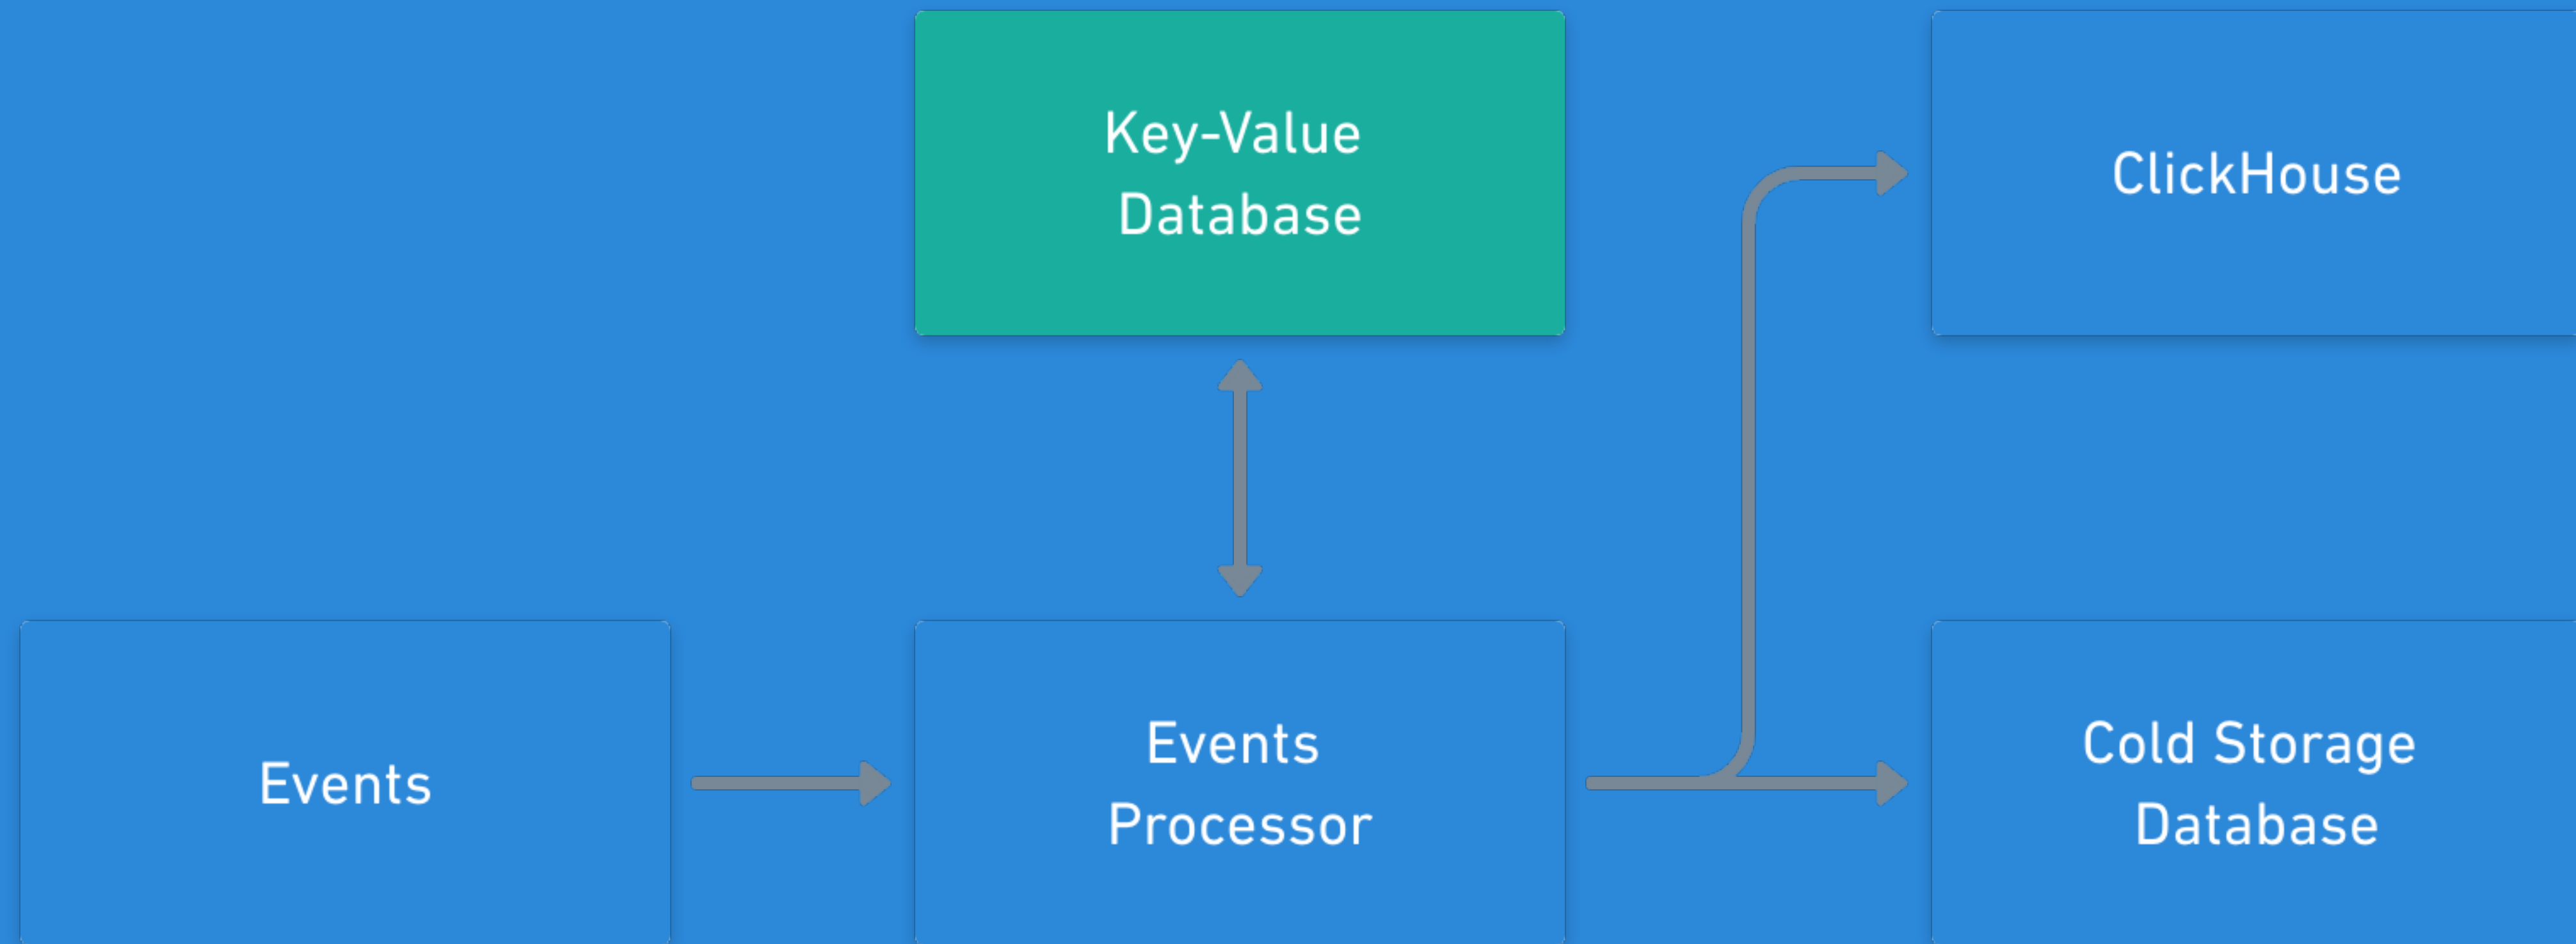




**How to insert the proper “negative” row?**

# **CollapsingMergeTree, keeping track of states**

- Need to be aware of the previous row to properly negate it
- ClickHouse is not made for random access of single rows





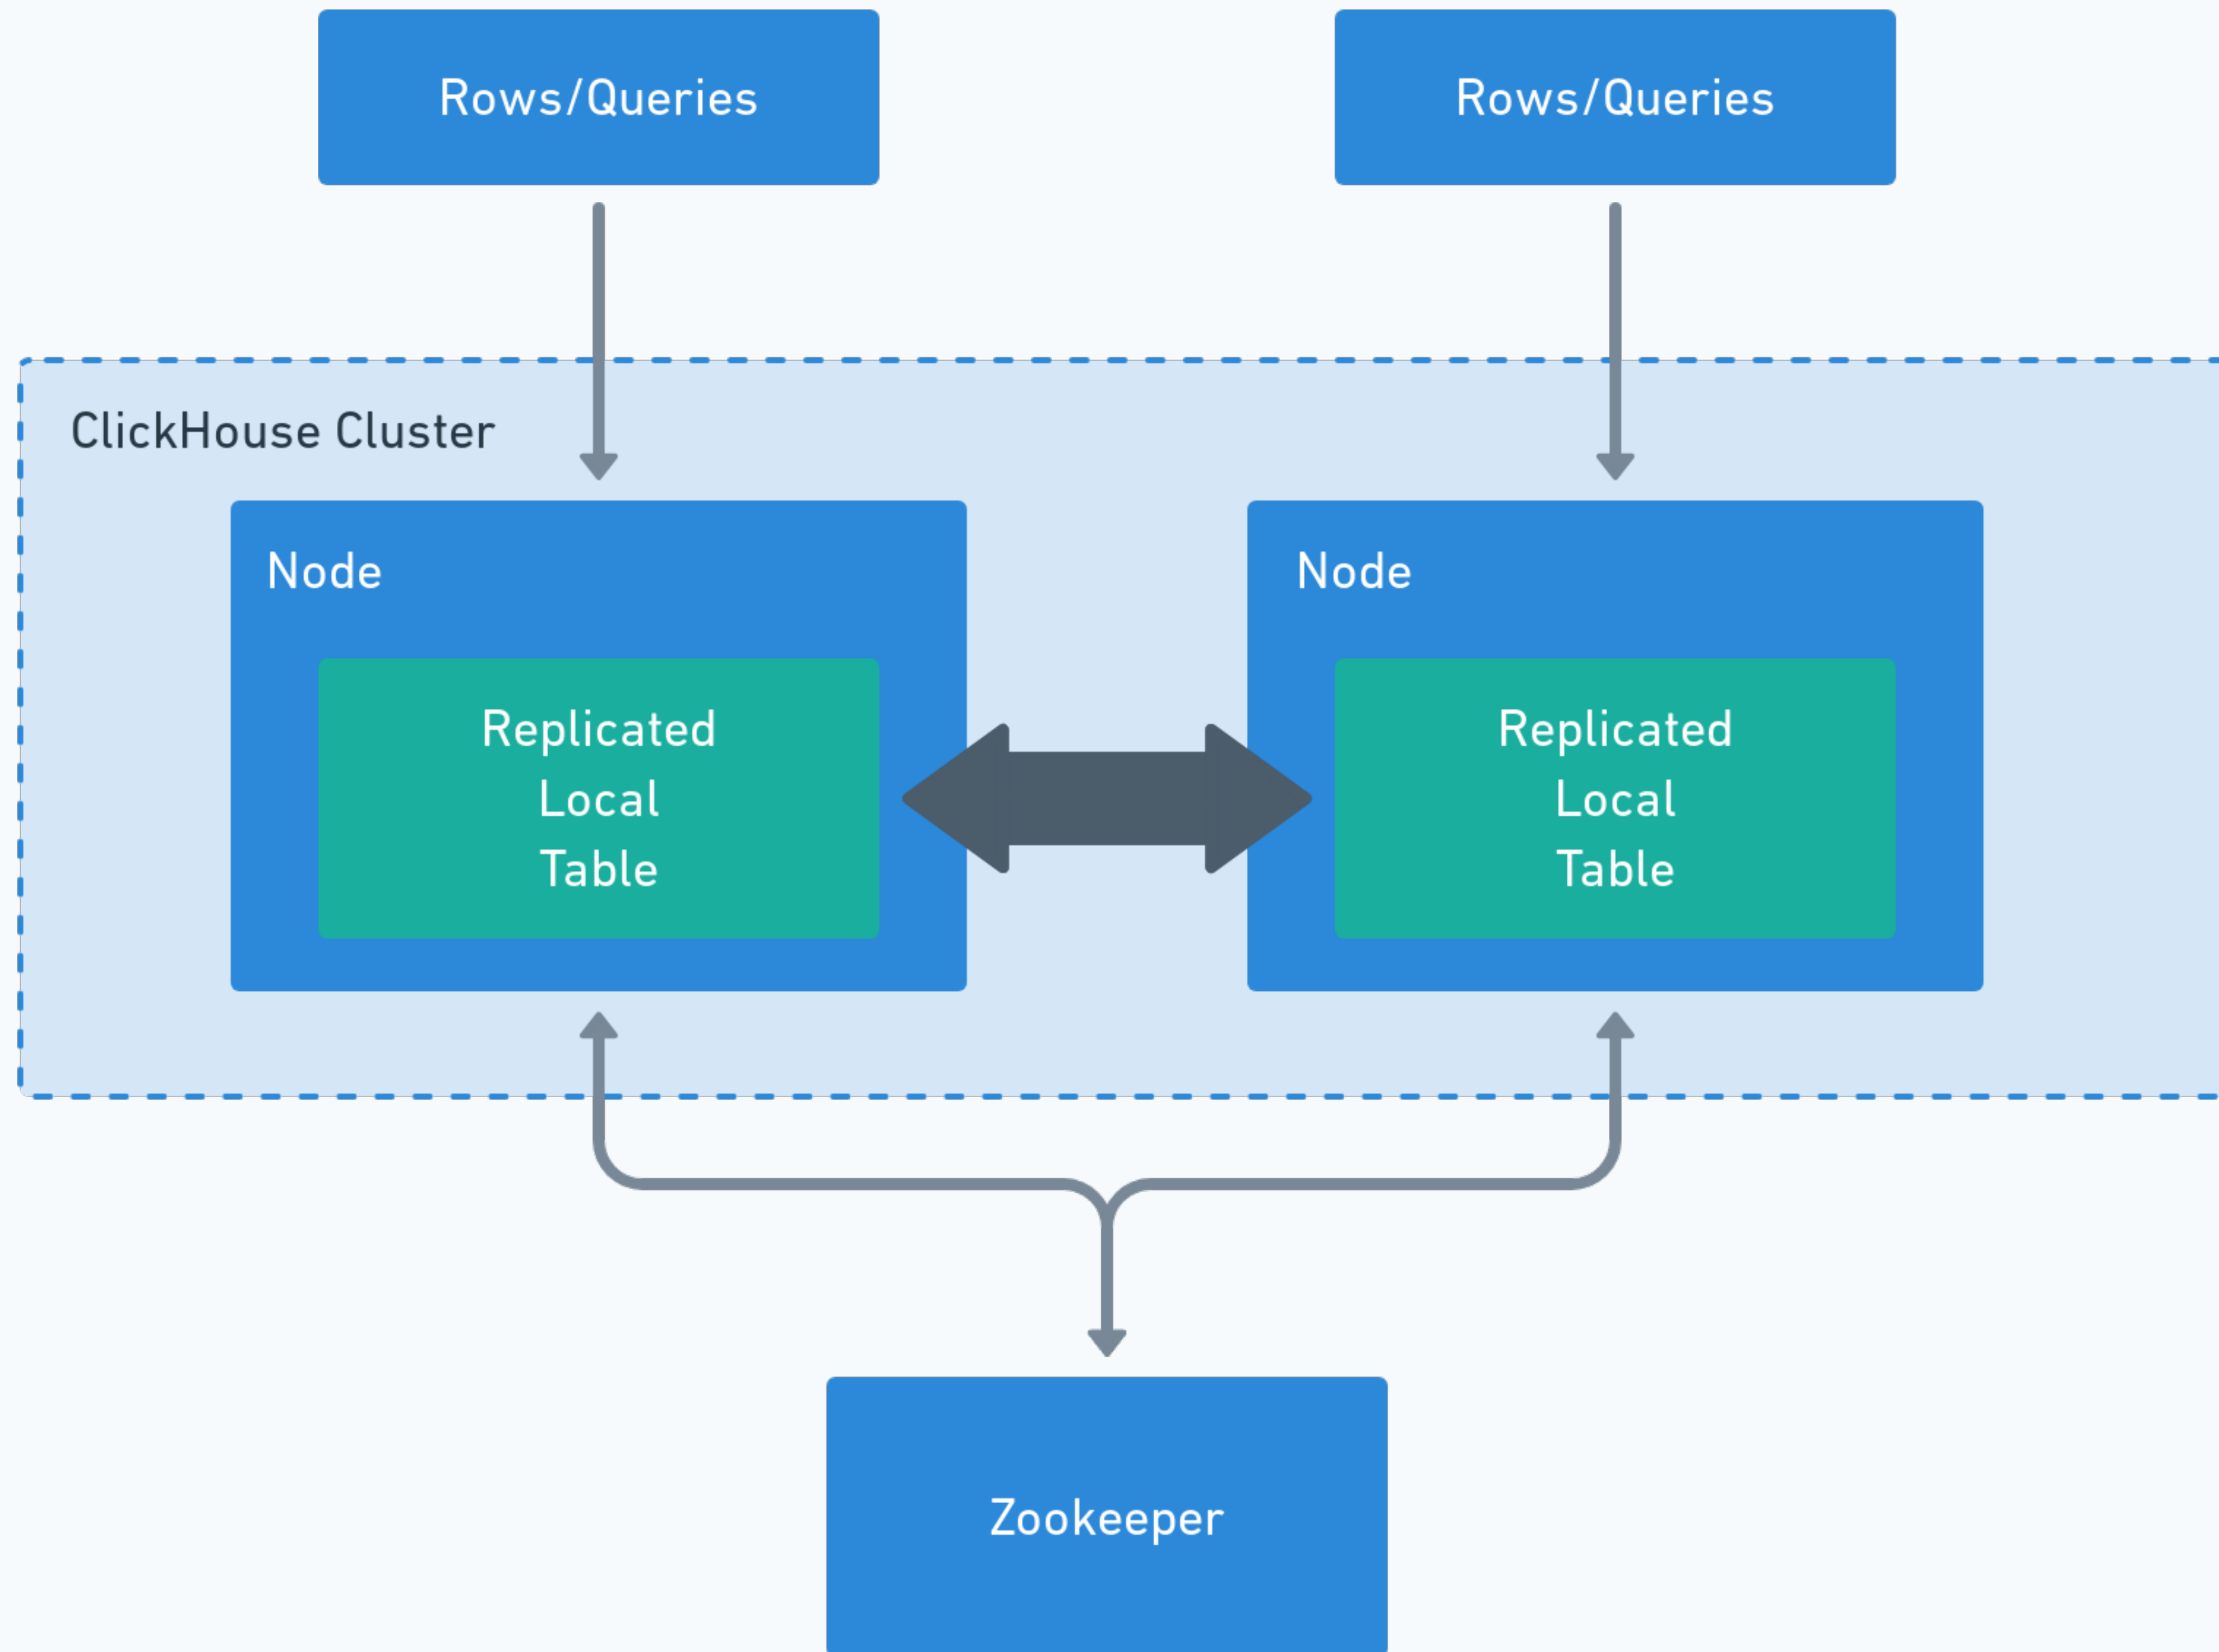
# What about availability?

# Replication

- High availability and reliability
- To bring data closer to consumer
- More than one way to do it with ClickHouse

# ReplicatedMergeTree\*

- Is supported by the MergeTree table family
  - ReplicatedCollapsingMergeTree
  - ReplicatedAggregatingMergeTree
- Uses Zookeeper to coordinate the replication between nodes



# ClickHouse scalability?





# Horizontal scalability

- Distributed engine
  - Dispatch read queries to all the nodes
  - Shard the data and dispatch it to the right node
- Flexible sharding capabilities
  - Let ClickHouse do the work
  - Shard manually: inserting directly into the wanted node and only use the distributed engine to dispatch read queries

# Vertical scalability

- Very efficient use of available CPU
- Data is on one machine (or even in memory) makes queries even faster
- You don't care about sharding of the data, operations can be done on local table
- Generally accepted to have more CPU, rather than more servers

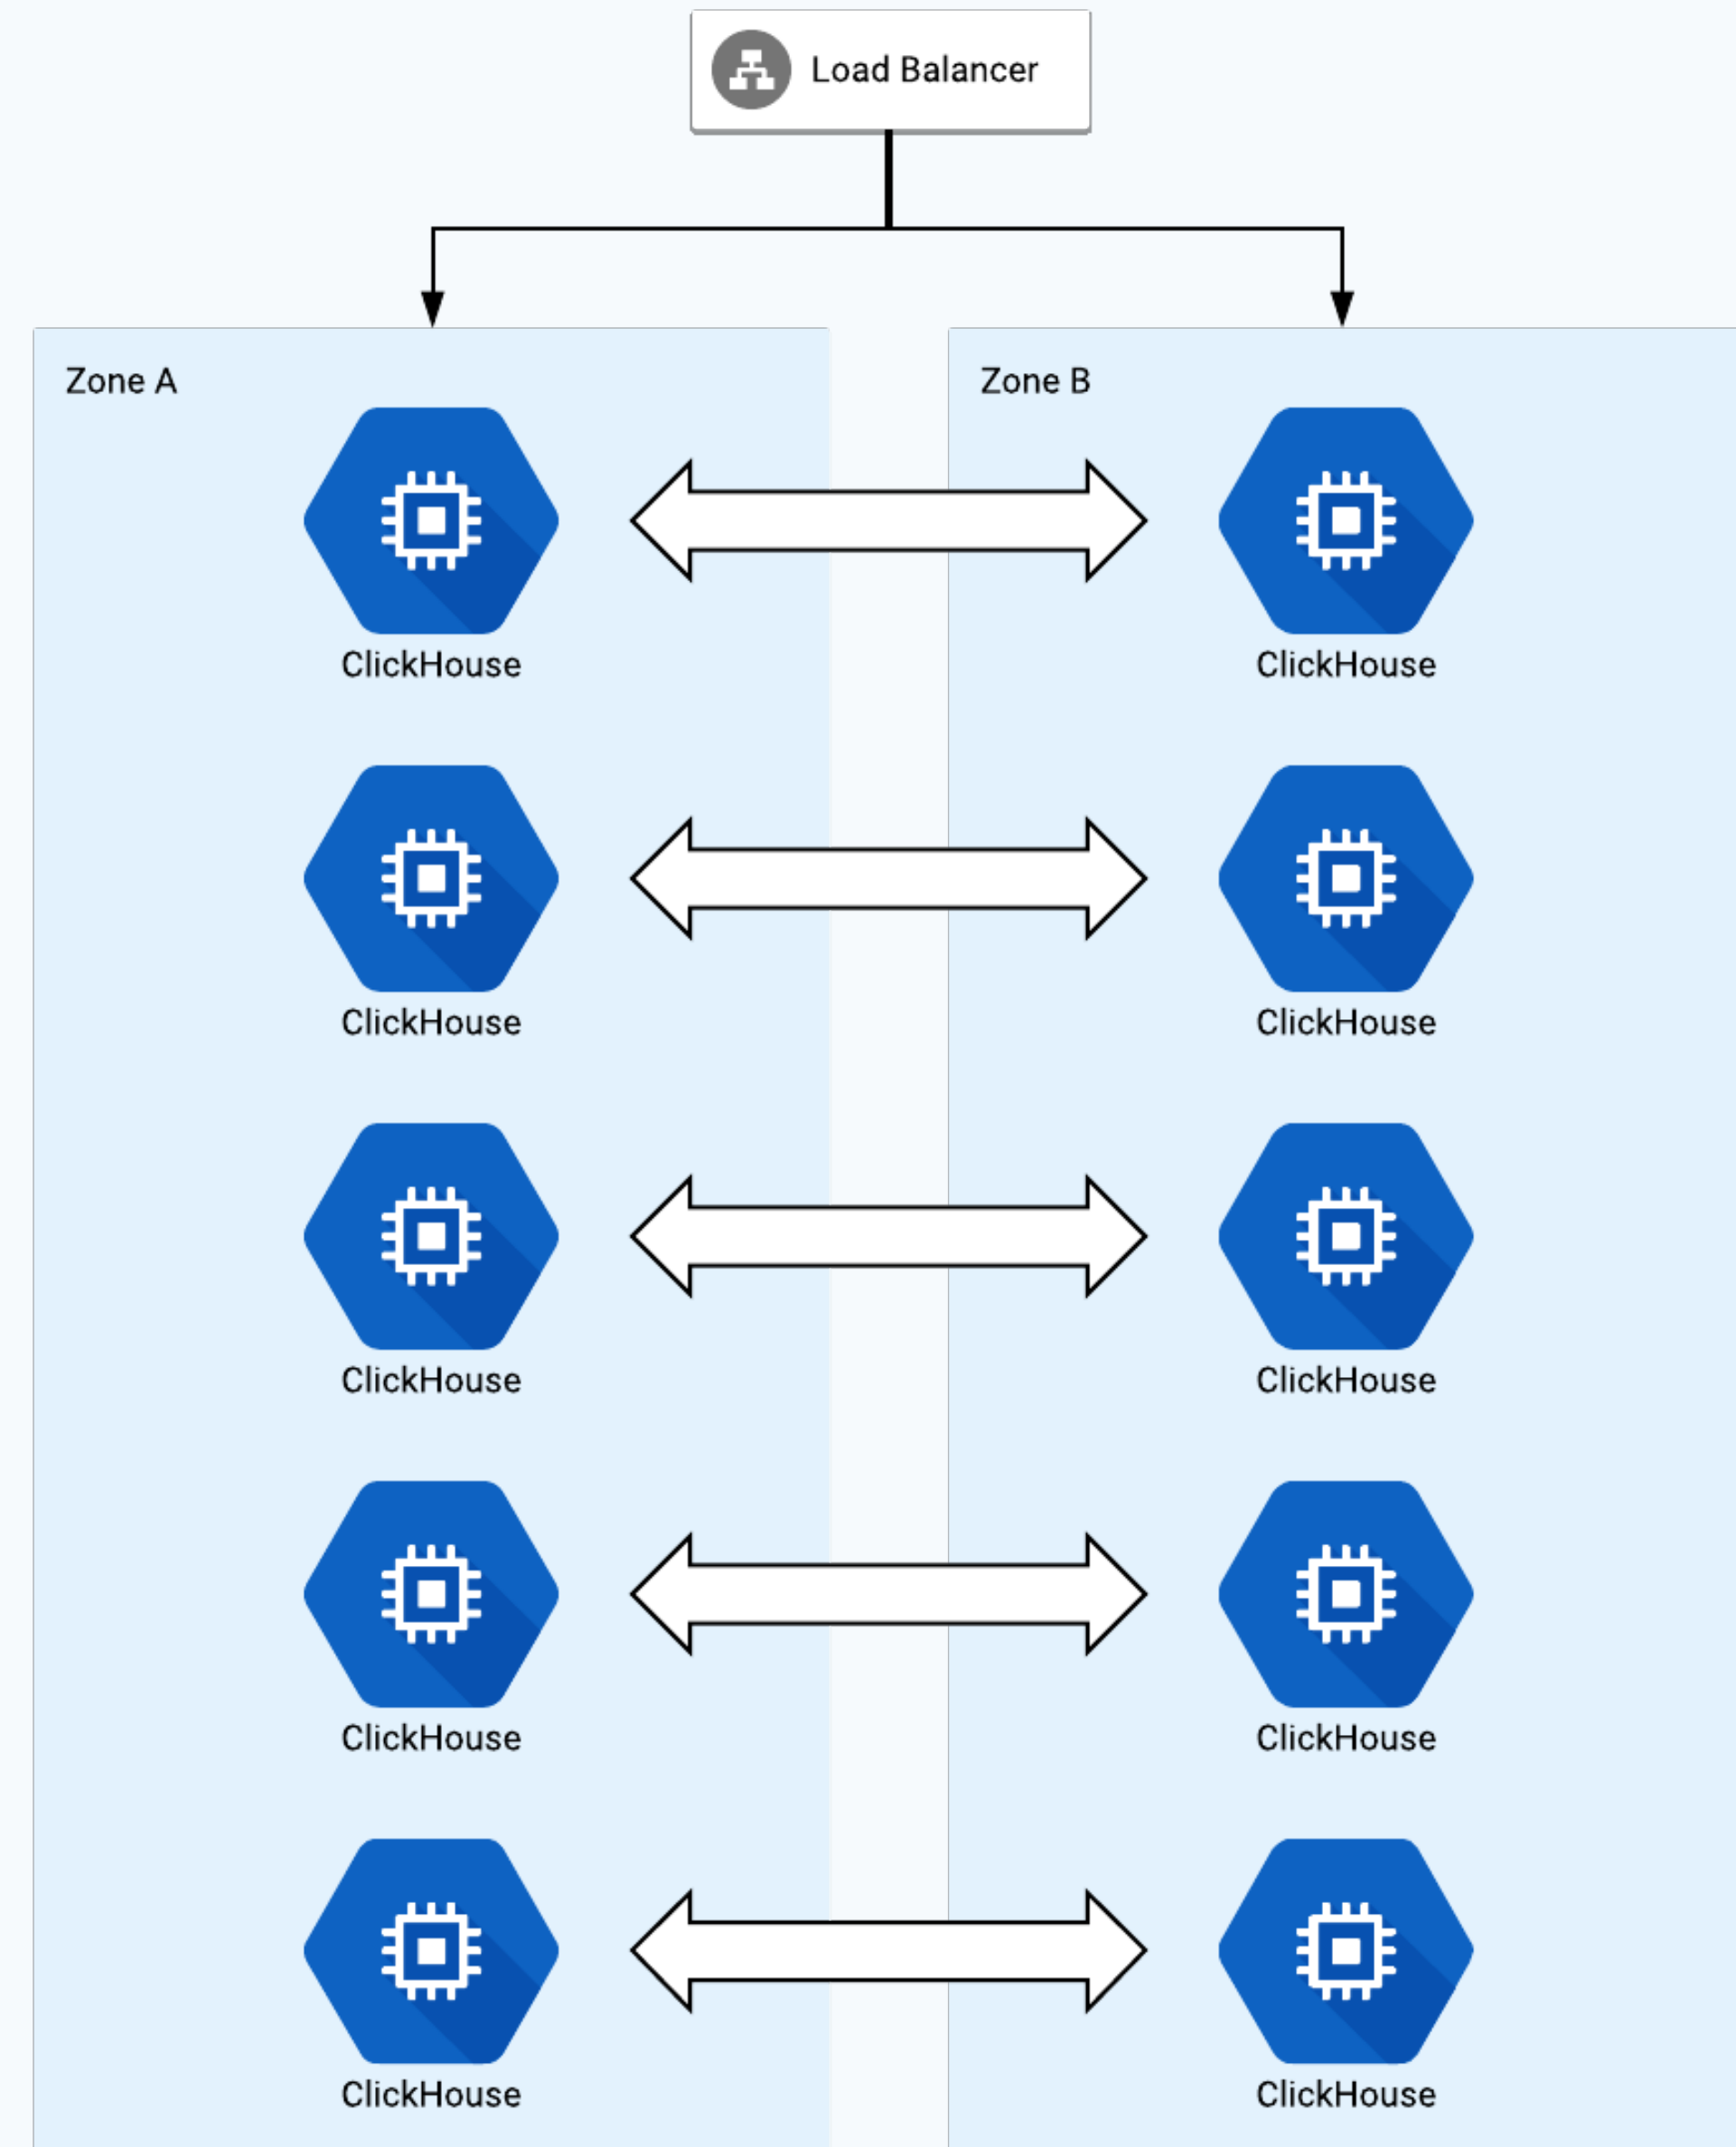
# ClickHouse in production?



## ENLIGHTENMENT - CLICKHOUSE USE CASE

# Our setup

- Single region
- Two availability zones
- 8 CPU/30 GB RAM
- 2TB+ compressed
- 10 nodes
- Replica factor 2



# How far ClickHouse took us

- Between the moment we designed and implemented our first data pipeline with ClickHouse from an average of 1000 events/s to 10000+ event/s without having to scale the cluster.
- Most of MessageBird products' data is in ClickHouse



# Rough edges

# Don't forget it's not a RDBMS

- Eventual consistency \*
- No transactions
- A single non unique primary index
- Limited support of JOIN
- Experimental features are experimental FOR REAL that stuff will break
- Resharding isn't out-of-the-box
- Not made for deleting/updating random rows

# ClickHouse among many

- ClickHouse is still one among many
- Dictionaries: periodically refreshed view of external databases
- JDBC/OBDC drivers, remote/local file, custom executable
- Non standard SQL can make third party like business intelligence tools integration can be challenging



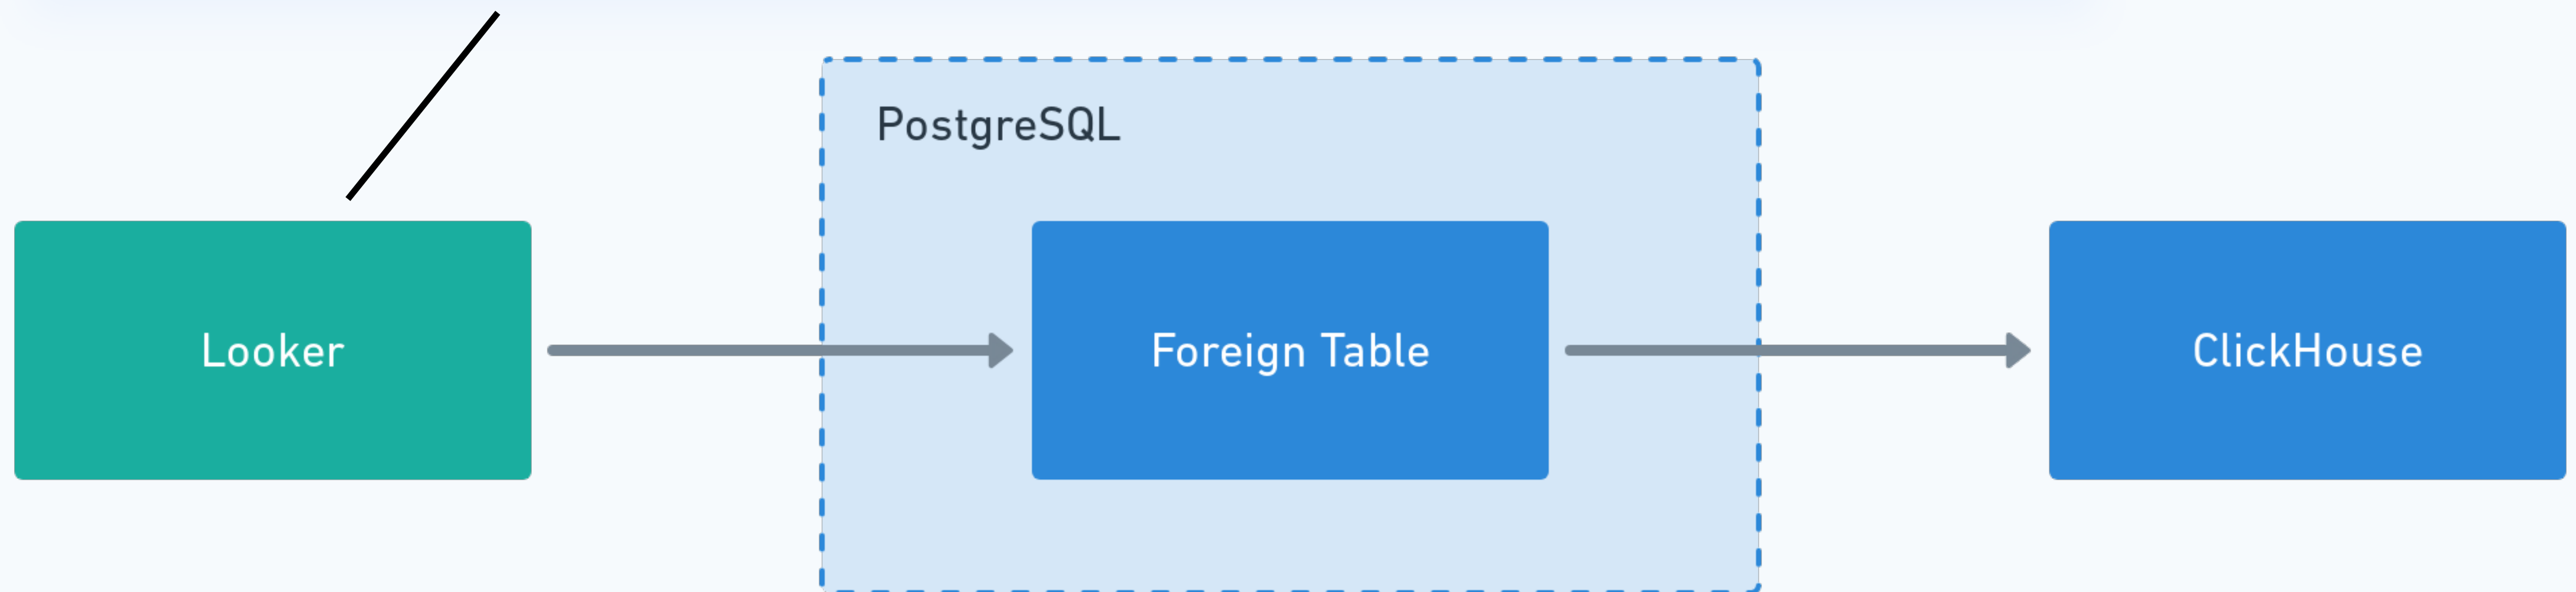
ENLIGHTENMENT - CLICKHOUSE USE CASE

# PostgreSQL + ClickHouse



# Query forwarding

```
SELECT sum(sign * price) AS total FROM dataset WHERE dataset = 666
```

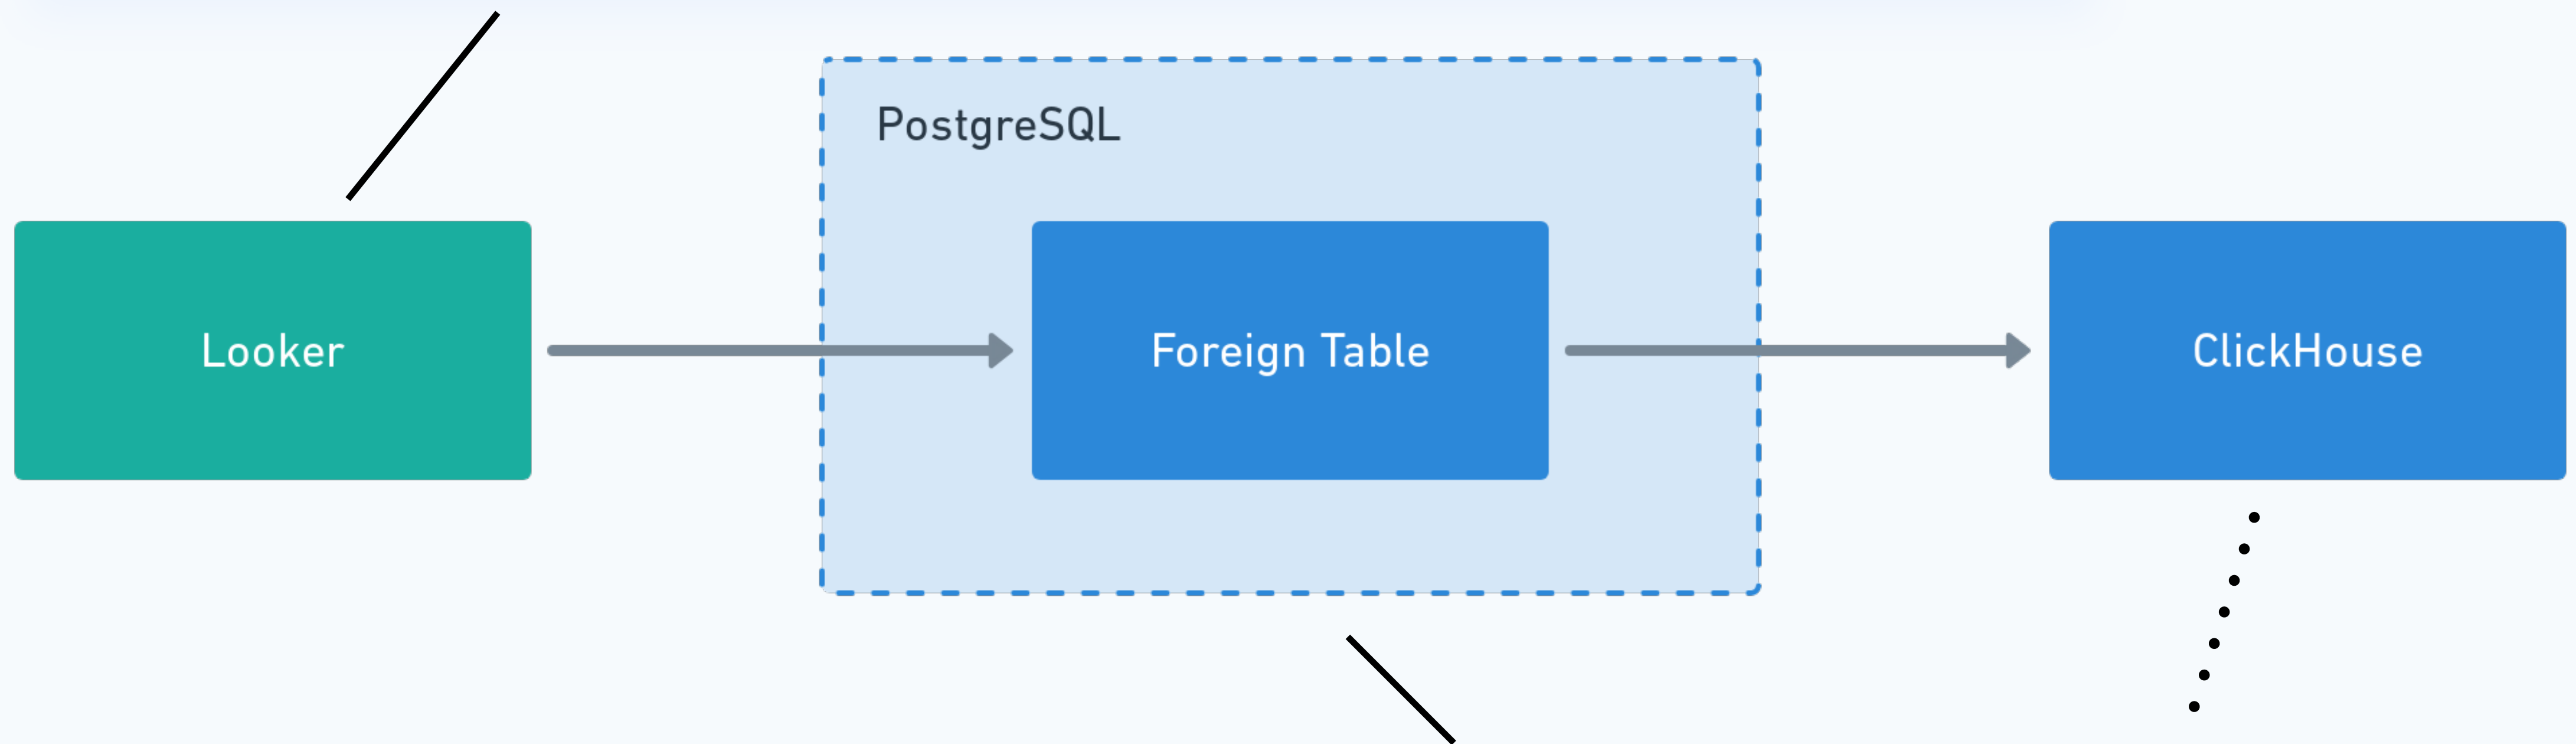


```
SELECT sign, price AS total FROM dataset
```



# Query forwarding and push down

```
SELECT sum(sign * price) AS total FROM dataset WHERE customer = 666
```

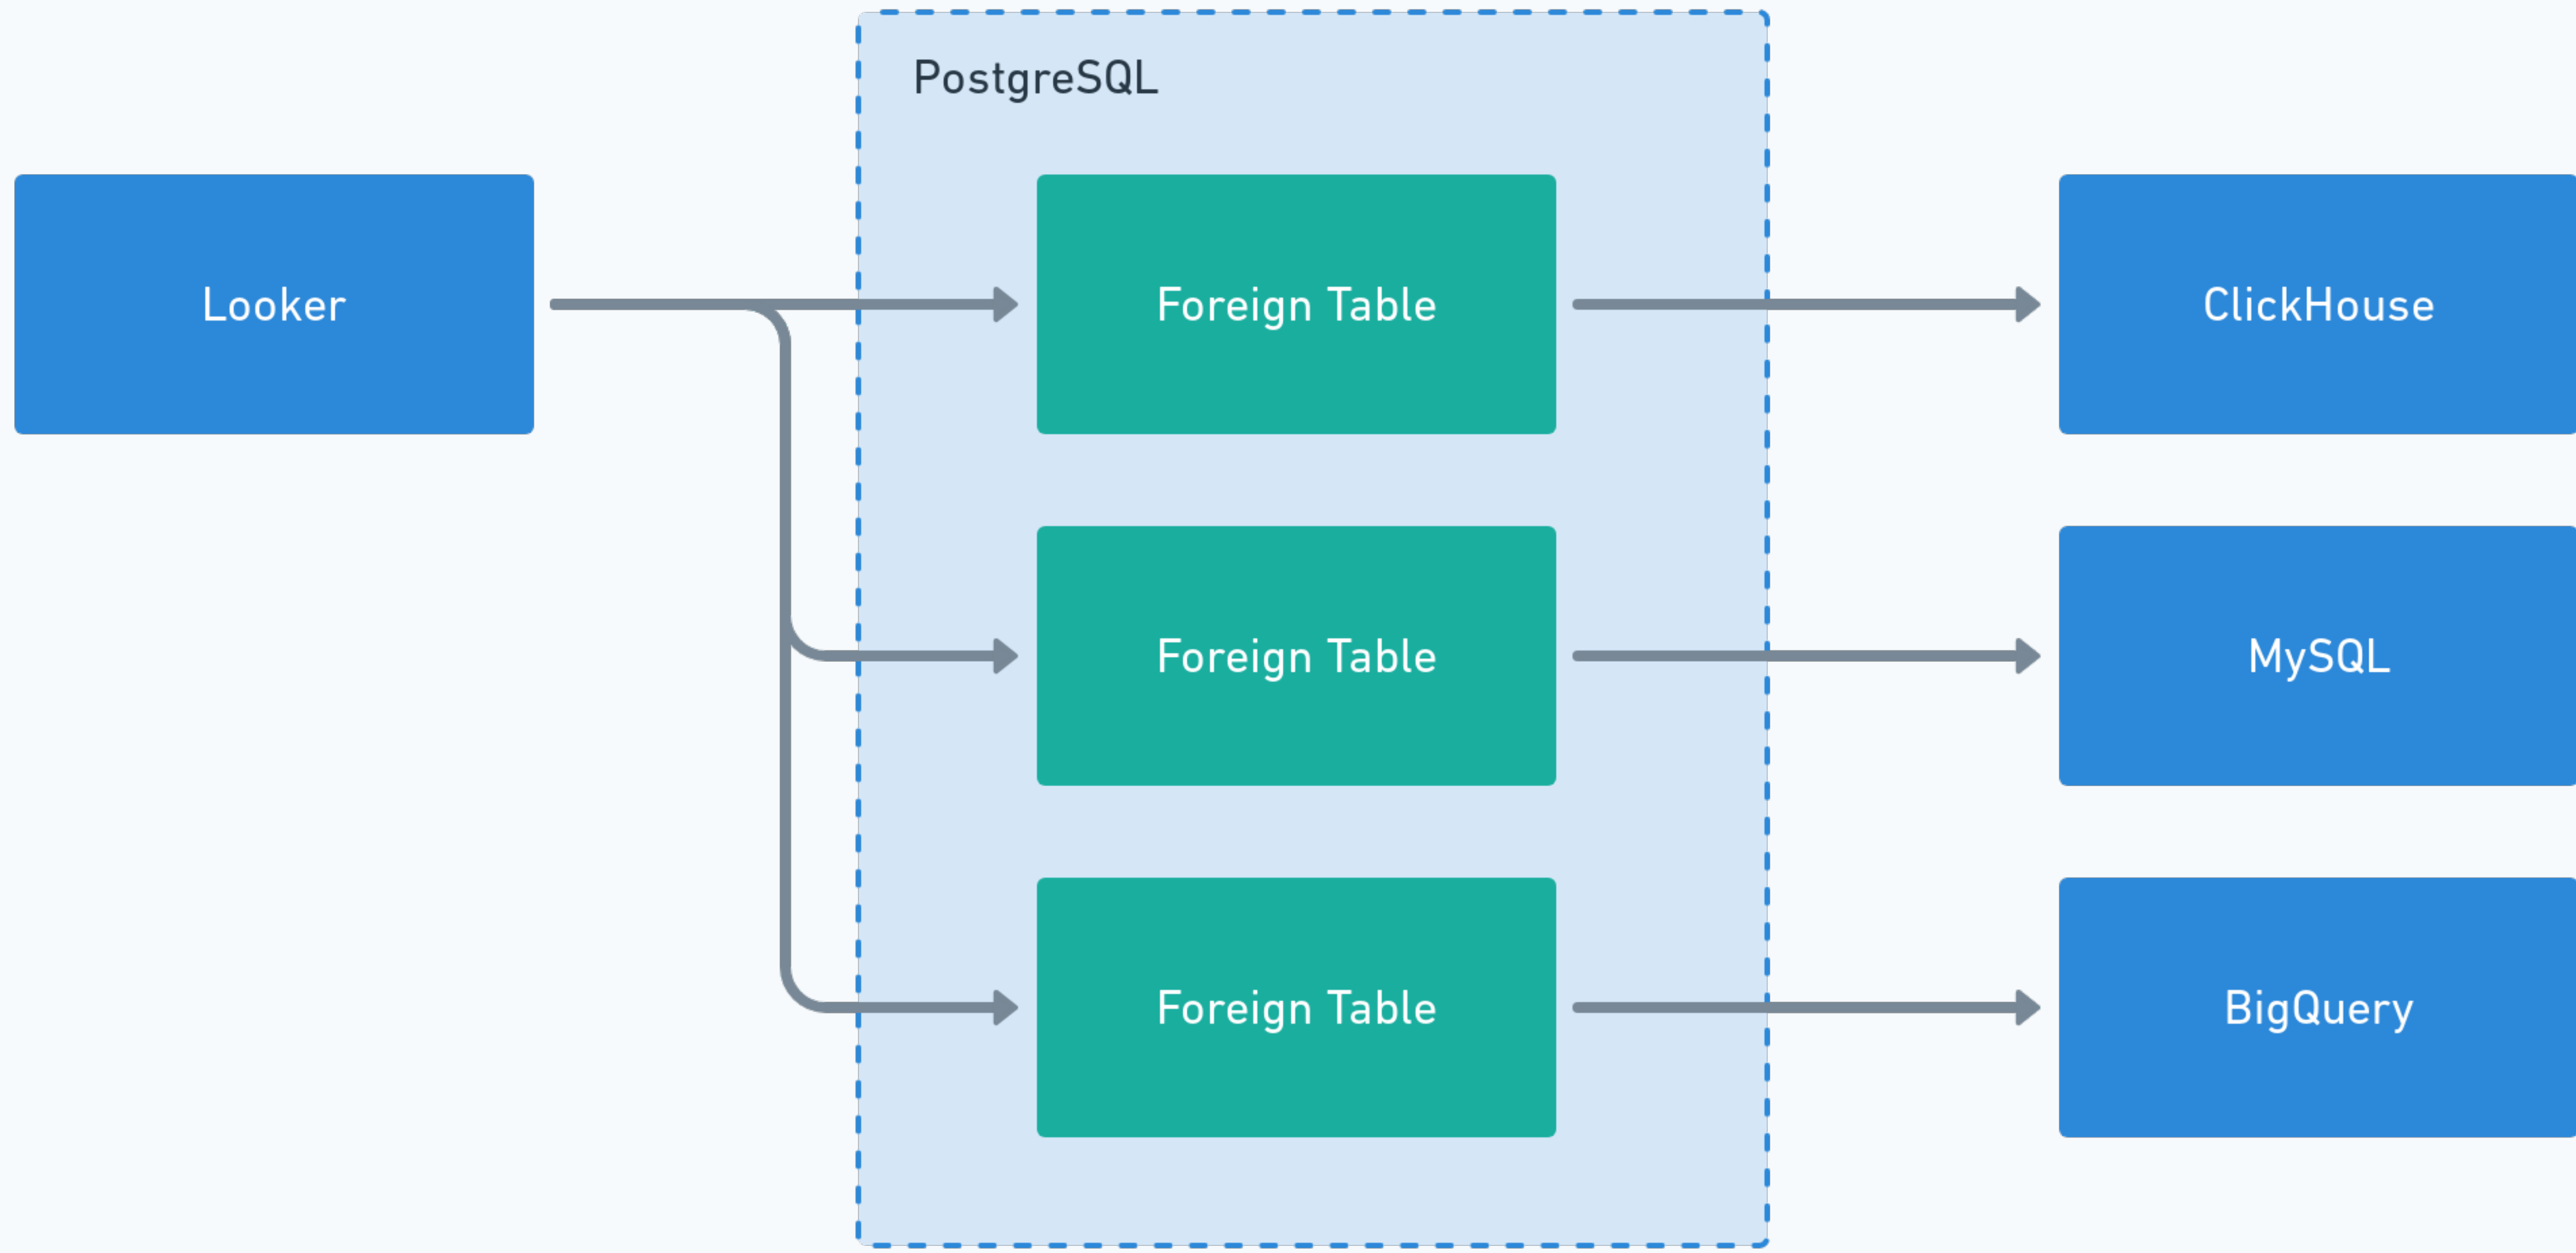


```
SELECT sign, price AS total FROM dataset WHERE customer = 666
```

# PostgreSQL + ClickHouse, looping the loop

- Instantly gain to one of the most standard SQL interface
- Still leverage the most important feature of ClickHouse by pushing down the filters and aggregations
- Bastion like approach to share data with third-party BI tools





ENLIGHTENMENT - CLICKHOUSE USE CASE

# PostgreSQL + ClickHouse, looping the loop

- Almost out-of-the-box data federation
- But only a PoC, we are still dreaming of production

*Did we say we are hiring?*



FUTURE - WRAPPING WORDS

# Even more possibilities

- ML features with catboost
- Kafka base table engine
- Upcoming better JOIN supports
- Cap'n Proto and upcoming Protobuf / Parquet support





# Questions

Late questions? Come say hello or drop us an email.

**[aleksandar@messagebird.com](mailto:aleksandar@messagebird.com) & [felix@messagebird.com](mailto:felix@messagebird.com)**

**[www.messagebird.com/careers](http://www.messagebird.com/careers)**



# Rate the session

